

```
import pandas as pd
```

```
df = pd.read_excel('/content/car.xlsx')
df
```



	brand	model	transmission	age	fuel	price	mileage	power	seats
0	18	244	1	4	1	1231000.0	19.01	4.496471	5
1	10	263	1	6	4	786000.0	19.01	4.496471	5
2	31	123	1	2	1	1489000.0	19.01	4.496471	5
3	9	55	0	1	4	1227000.0	19.01	4.496471	5
4	8	82	1	3	1	887000.0	19.01	4.496471	5
...
32009	5	199	1	6	4	292000.0	19.01	4.496471	5
32010	32	295	1	6	4	534000.0	19.01	4.496471	5
32011	33	25	1	8	4	424000.0	19.01	4.496471	5
32012	10	120	0	5	4	685000.0	19.01	4.496471	5
32013	31	247	1	2	4	392000.0	19.01	4.496471	5

32014 rows × 9 columns



Next steps:

[Generate code with df](#)
[View recommended plots](#)
[New interactive sheet](#)

```
X = df.drop('price',axis=1)
y = df['price']
```

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2,random_state=True)
```

```
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.model_selection import cross_val_score
```

```
# Initialize model
rf = RandomForestRegressor(random_state=42)
```

```
# Train
rf.fit(X_train, y_train)
```

```
# Predict
y_train_pred = rf.predict(X_train)
y_test_pred = rf.predict(X_test)
```

```
# Evaluate
print("Train R²:", r2_score(y_train, y_train_pred))
print("Test R²:", r2_score(y_test, y_test_pred))
print("Train MSE:", mean_squared_error(y_train, y_train_pred))
print("Test MSE:", mean_squared_error(y_test, y_test_pred))
```

```
# Cross-validation
cv_scores = cross_val_score(rf, X, y, cv=5, scoring='r2')
```

```
print("Cross-validation R2 scores:", cv_scores)
print("Mean CV R2 score:", cv_scores.mean())
```

↩ Train R²: 0.9558684229720562
 Test R²: 0.8918851685022822
 Train MSE: 22254754017.95725
 Test MSE: 49155901439.30192
 Cross-validation R² scores: [0.80217462 0.87467664 0.88807976 0.89361595 0.72405685]
 Mean CV R² score: 0.8365207636902584

```
from sklearn.model_selection import GridSearchCV
```

```
# Hyperparameter grid
```

```
param_grid = {
    'n_estimators': [1, 2],
    'max_depth': [None, 10, 20],
    'min_samples_split': [2, 5],
    'min_samples_leaf': [1, 2]
}
```

```
# GridSearchCV
```

```
grid = GridSearchCV(RandomForestRegressor(random_state=42), param_grid, cv=5,
                    scoring='r2', verbose=1, n_jobs=-1)
```

```
# Train
```

```
grid.fit(X_train, y_train)
```

```
# Best model
```

```
best_rf = grid.best_estimator_
```

```
# Predict
```

```
y_train_pred = best_rf.predict(X_train)
```

```
y_test_pred = best_rf.predict(X_test)
```

```
# Evaluate
```

```
print("🔑 Best Parameters:", grid.best_params_)
print("Train R2:", r2_score(y_train, y_train_pred))
print("Test R2:", r2_score(y_test, y_test_pred))
print("Train MSE:", mean_squared_error(y_train, y_train_pred))
print("Test MSE:", mean_squared_error(y_test, y_test_pred))
```

```
# Cross-validation on best model
```

```
cv_scores = cross_val_score(best_rf, X, y, cv=5, scoring='r2')
print("Cross-validation R2 scores:", cv_scores)
print("Mean CV R2 score:", cv_scores.mean())
```

↩ Fitting 5 folds for each of 24 candidates, totalling 120 fits
 🔑 Best Parameters: {'max_depth': 20, 'min_samples_leaf': 2, 'min_samples_split': 5, 'n_estimators': 200}
 Train R²: 0.9206326597464329
 Test R²: 0.879101461458048
 Train MSE: 40023510451.12758
 Test MSE: 54968190417.51255
 Cross-validation R² scores: [0.76552368 0.8454939 0.8666347 0.86270514 0.68885915]
 Mean CV R² score: 0.8058433142099147

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.